


Fraud & AML Detection System - Java + AI Architecture Guide

System Overview

A modern fraud and AML system that combines traditional rule-based detection with AI/ML models for intelligent pattern recognition.



Architecture Components

1. Data Ingestion Layer

- Real-time transaction streaming
- Batch customer data processing
- External data enrichment

2. AI/ML Detection Layer

- Anomaly detection
- Pattern recognition
- Risk scoring
- Behavioral analysis

3. Rules Engine

- Regulatory compliance (FATF, FinCEN)
- Threshold-based alerts
- Watchlist screening
- Sanctions checking

4. Case Management

- Alert investigation
- Suspicious Activity Reports (SARs)
- Workflow automation

5. Analytics & Reporting

- Real-time dashboards
 - Compliance reporting
 - Model performance monitoring
-

Technology Stack Recommendation

Core Java Stack

- Spring Boot 3.x (Application framework)
- Spring Cloud (Microservices)
- Apache Kafka (Event streaming)
- PostgreSQL (Transactional data)
- Redis (Caching, real-time scoring)
- Elasticsearch (Search, analytics)
- Apache Flink (Stream processing)
- Docker + Kubernetes (Containerization)


AI/ML Integration


- Python (ML Models) (Scikit-learn, TensorFlow, PyTorch)
- MLflow (Model management)
- FastAPI (ML model serving)
- Seldon Core / KServe (Model deployment on K8s)
- Apache Spark (Big data ML training)
- H2O.ai / Databricks (AutoML platforms)

Alternative: Java-Native ML

- Deeplearning4j (Deep learning in Java)
- Tribuo (Oracle) (ML library)
- Weka (Classic ML)
- Apache Mahout (Distributed ML)

System Architecture Diagram





Core Java Services Implementation

1. Transaction Monitoring Service

```
java
```

```
// TransactionMonitoringService.java
package com.fraudaml.monitoring;

import org.springframework.stereotype.Service;
import org.springframework.kafka.annotation.KafkaListener;
import lombok.RequiredArgsConstructor;
import lombok.extern.slf4j.Slf4j;

@Service
@Slf4j
@RequiredArgsConstructor
public class TransactionMonitoringService {

    private final RulesEngineService rulesEngine;
    private final MLScoreService mlScoreService;
    private final AlertService alertService;
    private final EnrichmentService enrichmentService;

    @KafkaListener(topics = "transactions", groupId = "fraud-detection")
    public void processTransaction(Transaction transaction) {
        log.info("Processing transaction: {}", transaction.getId());

        try {
            // 1. Enrich transaction data
            EnrichedTransaction enriched = enrichmentService.enrich(transaction);

            // 2. Get AI/ML risk score (async)
            CompletableFuture<MLScore> mlScore =
                mlScoreService.scoreTransactionAsync(enriched);

            // 3. Run rules engine
            RuleResult ruleResult = rulesEngine.evaluate(enriched);

            // 4. Combine ML score with rules
            MLScore score = mlScore.join(); // Wait for ML result

            // 5. Generate alert if needed
            if (shouldGenerateAlert(ruleResult, score)) {
                Alert alert = createAlert(enriched, ruleResult, score);
                alertService.publish(alert);
            }

            // 6. Store results
            storeAnalysisResult(enriched, ruleResult, score);
        } catch (Exception e) {
    }
}
```

```

        log.error("Error processing transaction: {}", transaction.getId(), e);
        // Dead letter queue handling
    }
}

private boolean shouldGenerateAlert(RuleResult ruleResult, MLScore mlScore) {
    // High ML score OR rule violations
    return mlScore.getScore() > 0.7 || ruleResult.hasViolations();
}

private Alert createAlert(EnrichedTransaction tx,
    RuleResult ruleResult,
    MLScore mlScore) {
    return Alert.builder()
        .transactionId(tx.getId())
        .alertType(determineAlertType(ruleResult, mlScore))
        .severity(calculateSeverity(mlScore.getScore()))
        .riskScore(mlScore.getScore())
        .violatedRules(ruleResult.getViolations())
        .features(mlScore.getFeatures())
        .timestamp(Instant.now())
        .status(AlertStatus.NEW)
        .build();
}
}

```

2. Transaction Entity

java

```
// Transaction.java
package com.fraudaml.domain;

import lombok.Data;
import java.math.BigDecimal;
import java.time.Instant;

@Data
public class Transaction {
    private String id;
    private String customerId;
    private String accountId;

    // Transaction details
    private BigDecimal amount;
    private String currency;
    private TransactionType type; // DEPOSIT, WITHDRAWAL, TRANSFER, PAYMENT
    private String channel; // ATM, ONLINE, MOBILE, BRANCH

    // Counterparty
    private String beneficiaryId;
    private String beneficiaryName;
    private String beneficiaryCountry;

    // Location
    private String ipAddress;
    private GeoLocation location;
    private String deviceId;

    // Timing
    private Instant timestamp;
    private String merchantCategory;

    // Additional
    private Map<String, Object> metadata;
}

@Data
class GeoLocation {
    private Double latitude;
    private Double longitude;
    private String country;
    private String city;
}
```

3. Rules Engine (Drools Integration)

```
java

// RulesEngineService.java
package com.fraudaml.rules;

import org.kie.api.runtime.KieSession;
import org.springframework.stereotype.Service;
import lombok.RequiredArgsConstructor;

@Service
@RequiredArgsConstructor
public class RulesEngineService {

    private final KieSession kieSession;

    public RuleResult evaluate(EnrichedTransaction transaction) {
        RuleResult result = new RuleResult();

        // Insert facts into working memory
        kieSession.insert(transaction);
        kieSession.insert(result);

        // Fire all rules
        kieSession.fireAllRules();

        // Dispose session
        kieSession.dispose();

        return result;
    }
}
```

4. Sample Drools Rules

```
java
```

```

// fraud-rules.drl

package com.fraudaml.rules;

import com.fraudaml.domain.EnrichedTransaction;
import com.fraudaml.rules.RuleResult;
import java.math.BigDecimal;

// Rule 1: High Value Transaction
rule "High Value Transaction"
when
    $tx : EnrichedTransaction(amount > 10000)
then
    RuleResult result = (RuleResult) kcontext.getKieRuntime()
        .getGlobal("result");
    result.addViolation("HIGH_VALUE",
        "Transaction exceeds $10,000: " + $tx.getAmount());
end

// Rule 2: Rapid Fire Transactions
rule "Rapid Fire Transactions"
when
    $tx : EnrichedTransaction()
    $count : Number(intValue >= 5) from accumulate(
        EnrichedTransaction(
            customerId == $tx.customerId,
            timestamp > $tx.timestamp.minusMinutes(5)
        ),
        count(1)
    )
then
    RuleResult result = (RuleResult) kcontext.getKieRuntime()
        .getGlobal("result");
    result.addViolation("RAPID_FIRE",
        "5+ transactions in 5 minutes");
end

// Rule 3: Unusual Location
rule "Unusual Location"
when
    $tx : EnrichedTransaction(
        location.country != customerProfile.homeCountry,
        timeSinceLastTransaction < 2 hours
    )
then
    RuleResult result = (RuleResult) kcontext.getKieRuntime()
        .getGlobal("result");

```

```

result.addViolation("LOCATION_ANOMALY",
    "Transaction from unusual location");
end

// Rule 4: Structuring Detection (AML)
rule "Structuring Detection"
when
    $count : Number(intValue >= 3) from accumulate(
        Transaction(
            customerId == $customerId,
            amount > 9000 && amount < 10000,
            timestamp > today.minusDays(1)
        ),
        count(1)
    )
then
    result.addViolation("STRUCTURING",
        "Multiple transactions just below reporting threshold");
end

// Rule 5: Sanctioned Country
rule "Sanctioned Country"
when
    $tx : EnrichedTransaction(
        beneficiaryCountry in ("IR", "KP", "SY", "CU")
    )
then
    result.addViolation("SANCTIONS",
        "Transaction to sanctioned country: " +
        $tx.getBeneficiaryCountry());
end

```

AI/ML Integration - Python Service

ML Scoring Service (FastAPI)

python

```
# ml_scoring_service.py
from fastapi import FastAPI, HTTPException
from pydantic import BaseModel
import numpy as np
import joblib
from typing import Dict, List
import logging

app = FastAPI(title="Fraud Detection ML Service")
logger = logging.getLogger(__name__)

# Load trained models
fraud_model = joblib.load('models/fraud_detector_v1.pkl')
aml_model = joblib.load('models/aml_detector_v1.pkl')
scaler = joblib.load('models/scaler.pkl')

class Transaction(BaseModel):
    id: str
    amount: float
    customer_id: str
    transaction_type: str
    channel: str
    beneficiary_country: str
    hour_of_day: int
    day_of_week: int
    customer_age_days: int
    avg_transaction_amount_30d: float
    transaction_count_24h: int
    customer_risk_score: float
    ip_country: str
    device_fingerprint: str

class MLScore(BaseModel):
    transaction_id: str
    fraud_score: float
    aml_score: float
    combined_score: float
    features: Dict[str, float]
    model_version: str
    explanation: List[str]

@app.post("/score", response_model=MLScore)
async def score_transaction(transaction: Transaction):
    """
    Score a transaction for fraud and AML risk
    """
    # Load models and scaler
    # Scale transaction features
    # Predict fraud and AML scores
    # Calculate combined score
    # Return results as MLScore object
```

```
try:
    # Feature engineering
    features = extract_features(transaction)

    # Scale features
    X = scaler.transform([features])

    # Get predictions
    fraud_prob = fraud_model.predict_proba(X)[0][1]
    aml_prob = aml_model.predict_proba(X)[0][1]

    # Combined score (weighted average)
    combined_score = 0.6 * fraud_prob + 0.4 * aml_prob

    # Feature importance for explainability
    feature_importance = get_feature_importance(X, fraud_model)

    # Generate explanation
    explanation = generate_explanation(
        transaction,
        fraud_prob,
        aml_prob,
        feature_importance
    )

    return MLScore(
        transaction_id=transaction.id,
        fraud_score=round(fraud_prob, 4),
        aml_score=round(aml_prob, 4),
        combined_score=round(combined_score, 4),
        features=feature_importance,
        model_version="v1.0.0",
        explanation=explanation
    )

except Exception as e:
    logger.error(f"Error scoring transaction {transaction.id}: {e}")
    raise HTTPException(status_code=500, detail=str(e))

def extract_features(tx: Transaction) -> List[float]:
    """
    Extract and engineer features for ML model
    """
    features = [
        tx.amount,
        tx.hour_of_day,
        tx.day_of_week,
```

```

tx.customer_age_days,
tx.avg_transaction_amount_30d,
tx.transaction_count_24h,
tx.customer_risk_score,

# Derived features
tx.amount / (tx.avg_transaction_amount_30d + 1), # Amount ratio
1 if tx.hour_of_day < 6 or tx.hour_of_day > 22 else 0, # Unusual hours
1 if tx.beneficiary_country != tx.ip_country else 0, # Country mismatch

# One-hot encoded categorical features
1 if tx.transaction_type == "WITHDRAWAL" else 0,
1 if tx.transaction_type == "TRANSFER" else 0,
1 if tx.channel == "ONLINE" else 0,
1 if tx.channel == "ATM" else 0,
]

return features

```

```

def get_feature_importance(X, model) -> Dict[str, float]:
    """
    Get top feature contributions for explainability
    """
    feature_names = [
        "amount", "hour", "day_of_week", "customer_age",
        "avg_amount_30d", "tx_count_24h", "customer_risk",
        "amount_ratio", "unusual_hours", "country_mismatch"
    ]

    # Get feature importance from model
    if hasattr(model, 'feature_importances_'):
        importance = model.feature_importances_
    else:
        # For linear models
        importance = np.abs(model.coef_[0])

    # Normalize and create dict
    importance_norm = importance / np.sum(importance)

    return {
        name: round(float(importance_norm), 4)
        for name, imp in zip(feature_names, importance_norm)
    }

def generate_explanation(tx: Transaction,
                        fraud_score: float,
                        aml_score: float,

```

```

features: Dict[str, float]) -> List[str]:
"""

Generate human-readable explanation
"""

explanations = []

# High risk factors
if fraud_score > 0.7:
    explanations.append("⚠️ High fraud risk detected")

    if features.get('amount_ratio', 0) > 0.15:
        explanations.append("• Transaction amount significantly exceeds customer norm")

    if features.get('unusual_hours', 0) > 0.1:
        explanations.append("• Transaction at unusual time")

    if tx.transaction_count_24h > 10:
        explanations.append(f"• High transaction velocity: {tx.transaction_count_24h} in 24h")

if aml_score > 0.7:
    explanations.append("⚠️ High AML risk detected")

    if features.get('country_mismatch', 0) > 0.1:
        explanations.append("• IP location differs from beneficiary country")

return explanations if explanations else ["✓ No significant risk factors detected"]

@app.get("/health")
async def health_check():
    return {"status": "healthy", "models_loaded": True}

@app.get("/models/info")
async def model_info():
    return {
        "fraud_model": "RandomForestClassifier",
        "aml_model": "XGBoostClassifier",
        "version": "v1.0.0",
        "last_trained": "2024-01-15"
    }

```

Java ML Service Client

```
java
```

```
// MLScoreService.java (Java client)
package com.fraudaml.ml;

import org.springframework.stereotype.Service;
import org.springframework.web.reactive.function.client.WebClient;
import reactor.core.publisher.Mono;
import lombok.RequiredArgsConstructor;

@Service
@RequiredArgsConstructor
public class MLScoreService {

    private final WebClient mlServiceClient;

    public CompletableFuture<MLScore> scoreTransactionAsync(
        EnrichedTransaction transaction) {

        return mlServiceClient
            .post()
            .uri("/score")
            .bodyValue(convertToMLRequest(transaction))
            .retrieve()
            .bodyToMono(MLScore.class)
            .toFuture();
    }

    public MLScore scoreTransaction(EnrichedTransaction transaction) {
        return mlServiceClient
            .post()
            .uri("/score")
            .bodyValue(convertToMLRequest(transaction))
            .retrieve()
            .bodyToMono(MLScore.class)
            .block();
    }

    private MLTransactionRequest convertToMLRequest(EnrichedTransaction tx) {
        return MLTransactionRequest.builder()
            .id(tx.getId())
            .amount(tx.getAmount().doubleValue())
            .customerId(tx.getCustomerId())
            .transactionType(tx.getType().name())
            .channel(tx.getChannel())
            .beneficiaryCountry(tx.getBeneficiaryCountry())
            .hourOfDay(tx.getTimestamp().get(ChronoField.HOUR_OF_DAY))
            .dayOfWeek(tx.getTimestamp().get(ChronoField.DAY_OF_WEEK))
    }
}
```

```
.customerAgeDays(tx.getCustomerProfile().getAgeDays())
.avgTransactionAmount30d(tx.getCustomerProfile().getAvgAmount30d())
.transactionCount24h(tx.getCustomerProfile().getTransactionCount24h())
.customerRiskScore(tx.getCustomerProfile().getRiskScore())
.ipCountry(tx.getLocation().getCountry())
.deviceFingerprint(tx.getDeviceId())
.build();
}

}
```

Alert Management Service

```
java
```

```
// AlertService.java
package com.fraudaml.alert;

import org.springframework.stereotype.Service;
import org.springframework.kafka.core.KafkaTemplate;
import lombok.RequiredArgsConstructor;
import lombok.extern.slf4j.Slf4j;

@Service
@Slf4j
@RequiredArgsConstructor
public class AlertService {

    private final KafkaTemplate<String, Alert> kafkaTemplate;
    private final AlertRepository alertRepository;
    private final NotificationService notificationService;

    public void publish(Alert alert) {
        // 1. Save to database
        Alert saved = alertRepository.save(alert);

        // 2. Publish to Kafka for real-time processing
        kafkaTemplate.send("alerts", alert.getId(), saved);

        // 3. Send notifications based on severity
        if (alert.getSeverity() == Severity.CRITICAL) {
            notificationService.sendUrgentAlert(alert);
        }

        log.info("Alert published: {} - Severity: {}, Score: {}",
            alert.getId(),
            alert.getSeverity(),
            alert.getRiskScore());
    }

    public List<Alert> getAlertsForInvestigation(String analystId) {
        return alertRepository.findByStatusAndAssignedTo(
            AlertStatus.NEW,
            analystId
        );
    }

    public Alert investigateAlert(String alertId, Investigation investigation) {
        Alert alert = alertRepository.findById(alertId)
            .orElseThrow(() -> new AlertNotFoundException(alertId));
    }
}
```

```
    alert.setStatus(AlertStatus.INVESTIGATING);
    alert.setInvestigation(investigation);
    alert.setAssignedTo(investigation.getAnalystId());

    return alertRepository.save(alert);
}

public Alert closeAlert(String alertId, AlertResolution resolution) {
    Alert alert = alertRepository.findById(alertId)
        .orElseThrow(() -> new AlertNotFoundException(alertId));

    alert.setStatus(AlertStatus.CLOSED);
    alert.setResolution(resolution);
    alert.setClosedAt(Instant.now());

    // If SAR is needed, trigger filing workflow
    if (resolution.isFileSAR()) {
        sarFilingService.initiateSAR(alert);
    }

    return alertRepository.save(alert);
}
}
```

Case Management

```
java
```

```
// CaseManagementService.java
package com.fraudaml.cases;

import org.springframework.stereotype.Service;
import lombok.RequiredArgsConstructor;

@Service
@RequiredArgsConstructor
public class CaseManagementService {

    private final CaseRepository caseRepository;
    private final CustomerRepository customerRepository;
    private final TransactionRepository transactionRepository;

    public Case createCase(Alert alert) {
        Customer customer = customerRepository.findById(alert.getCustomerId())
            .orElseThrow();

        Case investigationCase = Case.builder()
            .customerId(customer.getId())
            .customerName(customer.getName())
            .alerts(List.of(alert))
            .priority(calculatePriority(alert))
            .status(CaseStatus.OPEN)
            .createdAt(Instant.now())
            .build();

        return caseRepository.save(investigationCase);
    }

    public CaseView getCaseDetails(String caseId) {
        Case case = caseRepository.findById(caseId)
            .orElseThrow(() -> new CaseNotFoundException(caseId));

        // Gather all related data
        List<Transaction> transactions = transactionRepository
            .findByCustomerId(case.getCustomerId());

        Customer customer = customerRepository
            .findById(case.getCustomerId())
            .orElseThrow();

        List<Alert> relatedAlerts = alertRepository
            .findByCustomerId(case.getCustomerId());

        // Build comprehensive view
    }
}
```

```

return CaseView.builder()
    .caseInfo(case)
    .customer(customer)
    .transactions(transactions)
    .alerts(relatedAlerts)
    .timeline(buildTimeline(case, transactions, relatedAlerts))
    .riskAssessment(assessRisk(customer, transactions))
    .build();
}

public void escalateCase(String caseId, String reason) {
    Case case = caseRepository.findById(caseId)
        .orElseThrow();

    case.setPriority(Priority.CRITICAL);
    case.setEscalated(true);
    case.setEscalationReason(reason);
    case.setEscalatedAt(Instant.now());

    caseRepository.save(case);

    // Notify senior analysts
    notificationService.notifyEscalation(case);
}
}

```

Key ML Models to Implement

1. Anomaly Detection

python

```
# Isolation Forest for anomaly detection
from sklearn.ensemble import IsolationForest

def train_anomaly_detector(normal_transactions):
    model = IsolationForest(
        contamination=0.01, # 1% expected anomalies
        random_state=42
    )
    model.fit(normal_transactions)
    return model

# Usage
anomaly_score = model.predict(transaction_features)
# -1 = anomaly, 1 = normal
```

2. Fraud Classification

```
python

# XGBoost for fraud detection
import xgboost as xgb

def train_fraud_classifier(X_train, y_train):
    model = xgb.XGBClassifier(
        max_depth=6,
        learning_rate=0.1,
        n_estimators=100,
        scale_pos_weight=50 # Handle imbalanced data
    )
    model.fit(X_train, y_train)
    return model
```

3. Network Analysis (AML)

```
python
```

```
# Graph neural network for money laundering detection
```

```
import networkx as nx
```

```
def detect_laundering_networks(transactions):
```

```
    # Build transaction graph
```

```
    G = nx.DiGraph()
```

```
    for tx in transactions:
```

```
        G.add_edge(
```

```
            tx.sender,
```

```
            tx.receiver,
```

```
            amount=tx.amount,
```

```
            timestamp=tx.timestamp
```

```
)
```

```
# Detect suspicious patterns
```

```
# 1. Circular flows
```

```
cycles = nx.simple_cycles(G)
```

```
# 2. Layering (many hops)
```

```
paths = nx.all_simple_paths(G, source, target, cutoff=10)
```

```
# 3. Rapid movement
```

```
velocity = calculate_money_velocity(G)
```

```
return suspicious_patterns
```

Database Schema

```
sql
```

```
-- PostgreSQL Schema

-- Transactions
CREATE TABLE transactions (
    id UUID PRIMARY KEY,
    customer_id VARCHAR(50) NOT NULL,
    account_id VARCHAR(50) NOT NULL,
    amount DECIMAL(15,2) NOT NULL,
    currency VARCHAR(3) NOT NULL,
    transaction_type VARCHAR(20) NOT NULL,
    channel VARCHAR(20) NOT NULL,
    beneficiary_id VARCHAR(50),
    beneficiary_name VARCHAR(200),
    beneficiary_country VARCHAR(2),
    ip_address INET,
    device_id VARCHAR(100),
    merchant_category VARCHAR(50),
    timestamp TIMESTAMP NOT NULL,
    created_at TIMESTAMP DEFAULT NOW(),
    INDEX idx_customer (customer_id),
    INDEX idx_timestamp (timestamp),
    INDEX idx_amount (amount)
);

-- Alerts
```

```
CREATE TABLE alerts (
    id UUID PRIMARY KEY,
    transaction_id UUID REFERENCES transactions(id),
    customer_id VARCHAR(50) NOT NULL,
    alert_type VARCHAR(50) NOT NULL,
    severity VARCHAR(20) NOT NULL,
    risk_score DECIMAL(5,4) NOT NULL,
    violated_rules JSONB,
    ml_features JSONB,
    status VARCHAR(20) NOT NULL,
    assigned_to VARCHAR(50),
    created_at TIMESTAMP DEFAULT NOW(),
    updated_at TIMESTAMP DEFAULT NOW(),
    closed_at TIMESTAMP,
    INDEX idx_status (status),
    INDEX idx_customer (customer_id),
    INDEX idx_risk_score (risk_score DESC)
);

-- Cases
```

```
CREATE TABLE cases (
```

```
    id UUID PRIMARY KEY,  
    customer_id VARCHAR(50) NOT NULL,  
    customer_name VARCHAR(200),  
    priority VARCHAR(20) NOT NULL,  
    status VARCHAR(20) NOT NULL,  
    escalated BOOLEAN DEFAULT FALSE,  
    escalation_reason TEXT,  
    resolution VARCHAR(20),  
    resolution_notes TEXT,  
    sar_filed BOOLEAN DEFAULT FALSE,  
    sar_id VARCHAR(50),  
    created_at TIMESTAMP DEFAULT NOW(),  
    closed_at TIMESTAMP,  
    INDEX idx_status (status),  
    INDEX idx_customer (customer_id)  
);
```

-- Customer Risk Profiles

```
CREATE TABLE customer_risk_profiles (  
    customer_id VARCHAR(50) PRIMARY KEY,  
    base_risk_score DECIMAL(5,4) NOT NULL,  
    transaction_count_30d INT,  
    avg_transaction_amount DECIMAL(15,2),  
    max_transaction_amount DECIMAL(15,2),  
    high_risk_countries JSONB,  
    kyc_status VARCHAR(20),  
    pep_status BOOLEAN DEFAULT FALSE,  
    sanction_hit BOOLEAN DEFAULT FALSE,  
    updated_at TIMESTAMP DEFAULT NOW()  
);
```

-- Model Performance

```
CREATE TABLE model_performance (  
    id SERIAL PRIMARY KEY,  
    model_name VARCHAR(50) NOT NULL,  
    model_version VARCHAR(20) NOT NULL,  
    precision DECIMAL(5,4),  
    recall DECIMAL(5,4),  
    f1_score DECIMAL(5,4),  
    auc_roc DECIMAL(5,4),  
    false_positive_rate DECIMAL(5,4),  
    evaluation_date TIMESTAMP DEFAULT NOW()  
);
```

Deployment Configuration

Docker Compose

yaml

```
# docker-compose.yml
version: '3.8'

services:
  # Java Services
  transaction-service:
    build: ./transaction-service
    ports:
      - "8081:8081"
    environment:
      - KAFKA_BOOTSTRAP_SERVERS=kafka:9092
      - POSTGRES_URL=jdbc:postgresql://postgres:5432/fraudaml
      - ML_SERVICE_URL=http://ml-service:8000
    depends_on:
      - kafka
      - postgres
      - ml-service

  alert-service:
    build: ./alert-service
    ports:
      - "8082:8082"
    environment:
      - KAFKA_BOOTSTRAP_SERVERS=kafka:9092
      - POSTGRES_URL=jdbc:postgresql://postgres:5432/fraudaml

  # Python ML Service
  ml-service:
    build: ./ml-service
    ports:
      - "8000:8000"
    volumes:
      - ./models:/app/models
    environment:
      - MODEL_PATH=/app/models

  # Infrastructure
  postgres:
    image: postgres:15
    ports:
      - "5432:5432"
    environment:
      - POSTGRES_DB=fraudaml
      - POSTGRES_USER=admin
      - POSTGRES_PASSWORD=secure_password
    volumes:
```

- postgres-data:/var/lib/postgresql/data

kafka:

image: confluentinc/cp-kafka:7.5.0

ports:

- "9092:9092"

environment:

- KAFKA_ZOOKEEPER_CONNECT=zookeeper:2181

- KAFKA_ADVERTISED_LISTENERS=PLAINTEXT://kafka:9092

zookeeper:

image: confluentinc/cp-zookeeper:7.5.0

environment:

- ZOOKEEPER_CLIENT_PORT=2181

redis:

image: redis:7-alpine

ports:

- "6379:6379"

elasticsearch:

image: docker.elastic.co/elasticsearch/elasticsearch:8.11.0

ports:

- "9200:9200"

environment:

- discovery.type=single-node

- xpack.security.enabled=false

volumes:

postgres-data:

Key Features to Implement

Phase 1: Core Detection

-  Real-time transaction monitoring
-  Rule-based detection
-  ML-based risk scoring
-  Alert generation
-  Basic case management

Phase 2: Advanced ML

-  Anomaly detection models

-  Network analysis (graph ML)
-  Behavioral profiling
-  Auto-retraining pipeline
-  Model explainability (SHAP)

Phase 3: AML Compliance

-  Watchlist screening
-  Sanctions checking
-  PEP (Politically Exposed Persons) detection
-  SAR filing automation
-  Regulatory reporting

Phase 4: Advanced Features

-  Entity resolution
 -  Social network analysis
 -  Real-time streaming ML
 -  Feedback loops
 -  A/B testing for models
-

Performance Considerations

Scalability Targets

- **Transaction throughput:** 10,000+ TPS
- **Alert latency:** < 500ms
- **ML inference:** < 100ms
- **Dashboard refresh:** Real-time (WebSocket)

Optimization Strategies

1. **Caching:** Redis for customer profiles, risk scores
 2. **Batch processing:** Apache Flink for aggregations
 3. **Async processing:** Kafka for decoupling
 4. **Database partitioning:** Time-series partitioning
 5. **ML optimization:** Model quantization, ONNX runtime
-

Compliance & Regulations

Key Standards

- **FATF** (Financial Action Task Force)
- **Bank Secrecy Act (BSA)**
- **FinCEN** reporting
- **GDPR** (for EU customers)
- **PCI DSS** (payment data)

Audit Trail

- Log every transaction analysis
 - Track all alert dispositions
 - Record analyst actions
 - Maintain model versions
 - Store decision rationale
-

Testing Strategy

java

```

// TransactionMonitoringServiceTest.java
@SpringBootTest
class TransactionMonitoringServiceTest {

    @Test
    void shouldDetectHighValueTransaction() {
        // Given
        Transaction tx = createTransaction(15000.00);

        // When
        RuleResult result = rulesEngine.evaluate(tx);

        // Then
        assertTrue(result.hasViolation("HIGH_VALUE"));
    }

    @Test
    void shouldDetectRapidFireTransactions() {
        // Given
        List<Transaction> transactions = create5TransactionsIn3Minutes();

        // When
        transactions.forEach(tx -> service.processTransaction(tx));

        // Then
        verify(alertService, times(1))
            .publish(argThat(alert ->
                alert.getViolatedRules().contains("RAPID_FIRE")));
    }

    @Test
    void mlServiceIntegration() {
        // Given
        Transaction tx = createSuspiciousTransaction();

        // When
        MLScore score = mlScoringService.scoreTransaction(tx);

        // Then
        assertThat(score.getFraudScore()).isGreaterThanOrEqualTo(0.7);
    }
}

```

Metrics to Track

```
java

// Prometheus metrics
@Timed(value = "transaction.processing.time")
@Counted(value = "alerts.generated")
public void processTransaction(Transaction tx) {
    // ... processing logic
}

// Key metrics:
// - Transactions processed per second
// - Alert generation rate
// - False positive rate
// - Average investigation time
// - Model accuracy over time
// - System latency (p50, p95, p99)
```

Next Steps

1. Start with MVP

- Basic transaction ingestion
- Simple rules engine
- Alert management UI

2. Add ML Gradually

- Start with pre-trained models
- Collect labeled data
- Train custom models

3. Iterate Based on Feedback

- Tune rules based on false positives
- Retrain models monthly
- Add new detection patterns

4. Scale Infrastructure

- Add more Kafka partitions
- Scale ML service horizontally
- Implement caching layers

Resources & References

- **AML Regulations:** FATF Recommendations
 - **ML Libraries:** Scikit-learn, XGBoost, TensorFlow
 - **Java Frameworks:** Spring Boot, Apache Kafka, Drools
 - **Model Explainability:** SHAP, LIME
 - **Graph Analysis:** Neo4j, NetworkX
 - **Testing Data:** Synthetic transaction generators
-

Good luck building your fraud and AML system! This is a complex but highly rewarding project. Start small, iterate quickly, and always prioritize compliance and accuracy! 